

Class 1: Modelling with seasonality and the frequency domain

Andrew Parnell
andrew.parnell@ucd.ie



Learning outcomes

- Understand how to fit seasonal models in forecast and JAGS
- Understand seasonal differencing and sARIMA models
- Know the difference between time and frequency domain models and be able to implement a basic Fourier model

1 / 27

2 / 27

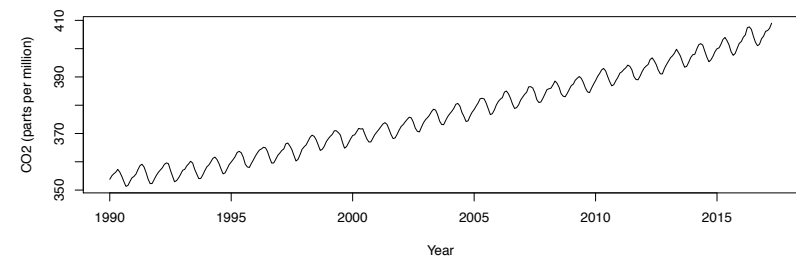
Seasonal time series

- So far we haven't covered how to deal with data that are *seasonal* in nature
- These data generally fall into two categories:
 1. Data where we know the frequency or frequencies (e.g. monthly data on a yearly cycle, frequency = 12)
 2. Data where we want to estimate the frequencies (e.g. climate time series, animal populations, etc)
- The former are easier, and there are many techniques for inducing seasonal behaviour
- The latter are much more interesting. The ACF and PACF can help, but we can usually do much better by creating a *power spectrum*

3 / 27

An example seasonal series

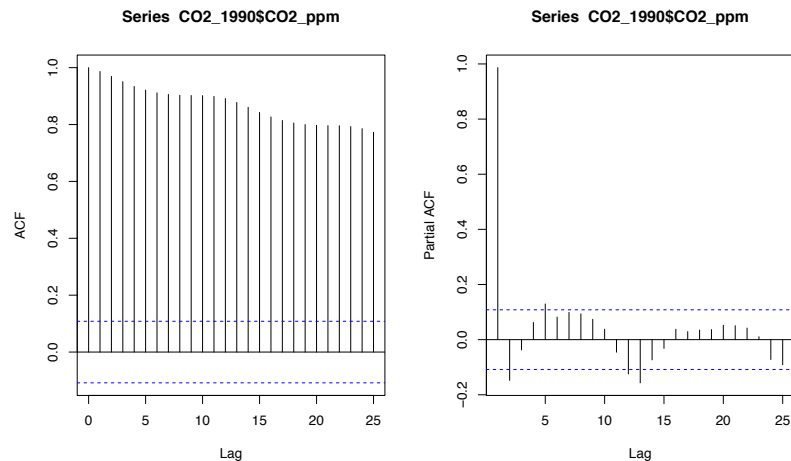
```
C02 = read.csv(file = '../data/C02.csv',  
               na.strings = -99.99)  
C02$date = as.Date(paste(C02$year, C02$month, '01',  
                         sep = '-'))  
C02_1990 = C02[C02$year >= 1990, ]  
with(C02_1990, plot(date, C02_ppm, type = 'l',  
                   ylab = 'CO2 (parts per million)',  
                   xlab = 'Year'), las = 1)
```



4 / 27

ACF and PACF

```
par(mfrow = c(1, 2))
acf(CO2_1990$CO2_ppm)
pacf(CO2_1990$CO2_ppm)
```



5 / 27

Seasonal time series 1: including seasonality as a covariate

- The simplest way is to include month as a covariate in a regression type model


```
CO2_1990$mfac = model.matrix(~ as.factor(CO2_1990$month) - 1)
colnames(CO2_1990$mfac) = month.abb
lm(CO2_ppm ~ year + mfac, data = CO2_1990)
```

```
##
## Call:
## lm(formula = CO2_ppm ~ year + mfac, data = CO2_1990)
##
## Coefficients:
## (Intercept)      year      mfacJan
## -3501.5897      1.9362     -0.6749
##      mfacFeb      mfacMar      mfacApr
##      0.1222      1.0272      2.3590
##      mfacMay      mfacJun      mfacJul
##      2.7967      2.1367      0.5126
##      mfacAug      mfacSep      mfacOct
##      -1.5704     -3.0774     -2.8770
```

6 / 27

Forecasts

```
CO2_ts = ts(CO2_1990$CO2_ppm, frequency = 12,
            start = c(1990, 1))
s_model_1 = tslm(CO2_ts ~ trend + season)
plot(forecast(s_model_1, h = 24))
```



7 / 27

What is the time series model doing here?

- This is just a regression model, so that:

$$y_t = \alpha + \beta_t \text{year}_t + \gamma_{1t} \text{Feb}_t + \gamma_{2t} \text{Mar}_t + \dots + \gamma_{11,t} \text{Dec}_t + \epsilon_t$$

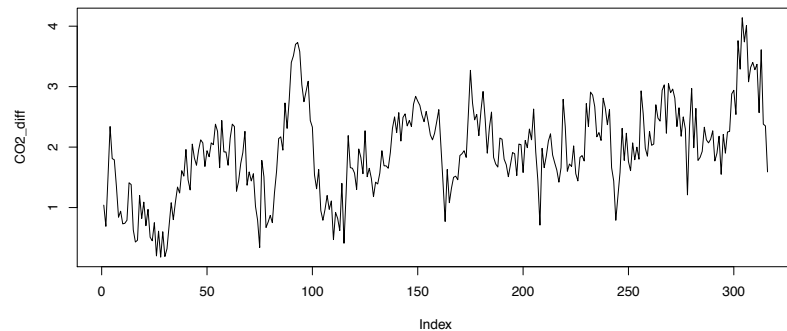
- You can do this using `lm` or using `forecast`'s special function for linear regression forecasting `tslm`
- The `tslm` function is clever because it can automatically create the seasonal indicator variables
- Remember that when dealing with indicator variables you have to drop one factor level for the model to fit

8 / 27

Seasonal time series 2: seasonal differencing

- ▶ We have already met methods which difference the data (possibly multiple times) at lag 1
- ▶ We can alternatively create a seasonal difference by differencing every e.g. 12th observation

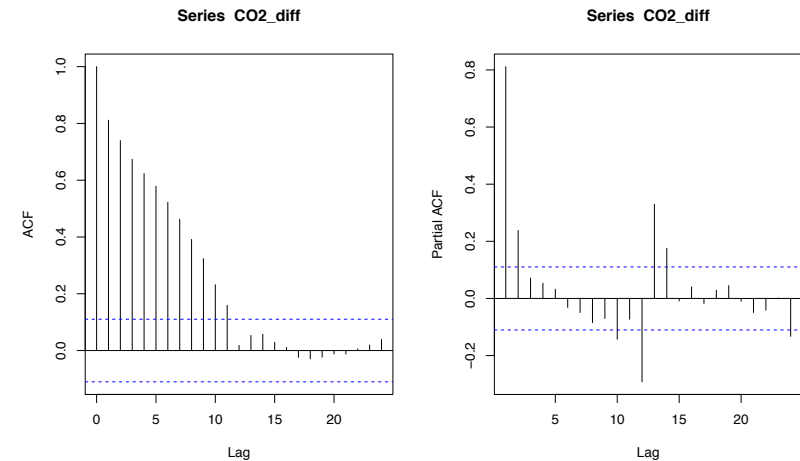
```
CO2_diff = diff(CO2_1990$CO2_ppm, lag = 12)
plot(CO2_diff, type = 'l')
```



9 / 27

Differenced acf and pacf

```
par(mfrow = c(1, 2))
acf(CO2_diff, na.action = na.pass)
pacf(CO2_diff, na.action = na.pass)
```



10 / 27

Fit an ARIMA model with a seasonal difference

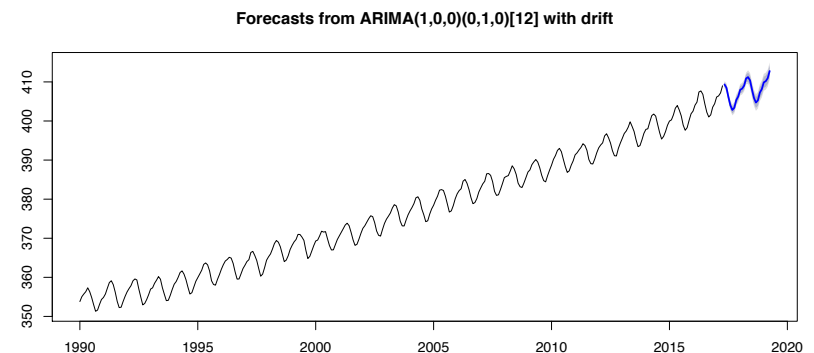
```
CO2_1990_ts = ts(CO2_1990$CO2_ppm, frequency = 12,
                 start = c(1990, 1))
Arima(CO2_1990_ts, order = c(1, 0, 0),
      seasonal = c(0, 1, 0),
      include.drift = TRUE)
```

```
## Series: CO2_1990_ts
## ARIMA(1,0,0)(0,1,0)[12] with drift
##
## Coefficients:
##      ar1      drift
##      0.8129  0.1588
## s.e.   0.0325  0.0107
##
## sigma^2 estimated as 0.1943:  log likelihood=-185.12
## AIC=376.23   AICc=376.31   BIC=387.5
```

11 / 27

Forecasts from seasonally differenced series

```
s_model_2 = Arima(CO2_1990_ts, order = c(1, 0, 0),
                  seasonal = c(0, 1, 0), include.drift = TRUE)
plot(forecast(s_model_2, h = 24))
```



- ▶ Pretty good. Might be able to do better with some richer models

12 / 27

A full seasonal arima model

- ▶ We previously met the ARIMA specification where:

$$\text{diff}^d(y_t) = \text{constant} + \text{AR terms} + \text{MA terms} + \text{error}$$

- ▶ We can extend this to include seasonal differencing and *seasonal AR and MA* terms to create a seasonal ARIMA or sARIMA model
- ▶ For example:

$$y_t - y_{t-12} = \alpha + \beta y_{t-1} + \gamma y_{t-12} + \epsilon_t$$

- ▶ This is a sARIMA(1,0,0)(1,1,0)₁₂ model

13 / 27

Fitting sARIMA models in forecast


```
auto.arima(CO2_1990_ts)
```

```
## Series: CO2_1990_ts
## ARIMA(0,1,1)(1,1,2)[12]
##
## Coefficients:
##          ma1      sar1      sma1      sma2
##      -0.3888  -0.7684  -0.1020  -0.6482
## s.e.   0.0582   0.4837   0.4891   0.4246
##
## sigma^2 estimated as 0.1137:  log likelihood=-103.27
## AIC=216.55   AICc=216.74   BIC=235.31
```

14 / 27

Plotting forecasts

```
s_model_3 = auto.arima(CO2_1990_ts)
plot(forecast(s_model_3, h = 24))
```



15 / 27

A simple sARIMA model with JAGS

```
model_code = '
model
{
  # Likelihood
  for (t in (s+1):T) {
    y[t] ~ dnorm(mu[t], sigma^-2)
    mu[t] <- alpha + beta * y[t-1] + gamma * y[t-s]
  }

  # Priors
  alpha ~ dnorm(0, 10^-2)
  beta ~ dnorm(0, 10^-2)
  gamma ~ dnorm(0, 10^-2)
  sigma ~ dunif(0, 100)
}
'
```

16 / 27

Fitting a sARIMA(1, 0, 0)(1, 0, 0)₁₂ model in JAGS

```
s_model_4 = jags(data = list(y = CO2_ts, s = 12,
                             T = length(CO2_ts)),
                 parameters.to.save = c('alpha', 'beta',
                                         'gamma', 'sigma'),
                 model.file=textConnection(model_code))

print(s_model_4)

## Inference for Bugs model at "5", fit using jags,
## 3 chains, each with 2000 iterations (first 1000 discarded)
## n.sims = 3000 iterations saved
##      mu.vect sd.vect   2.5%   25%   50%
## alpha  -6.387   0.873  -8.098  -6.969  -6.386
## beta    0.189   0.024   0.144   0.173   0.189
## gamma   0.832   0.024   0.783   0.816   0.832
## sigma   0.599   0.024   0.554   0.582   0.599
## deviance 571.705 2.890 568.132 569.615 571.055
##      75%   97.5%  Rhat n.eff
## alpha  -5.790  -4.689 1.001 3000
## beta    0.205   0.238 1.002 1600
## gamma   0.848   0.879 1.002 1600
## sigma   0.614   0.648 1.001 3000
## deviance 573.074 579.233 1.002 1800
##
## For each parameter, n.eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor (at convergence, Rhat=1).
##
## DIC info (using the rule, pD = var(deviance)/2)
## pD = 4.2 and DIC = 575.9
## DIC is an estimate of expected predictive error (lower deviance is better).
```

17 / 27

Frequency estimation

Multiple seasonality

- ▶ Very occasionally you come across multiple seasonality models
- ▶ For example you might have hourly data over several months with both hourly and monthly seasonality
- ▶ forecast has a special function for creating multiple series time series: `msts`

```
x = msts(taylor, seasonal.periods=c(48,336),
        start=2000+22/52)
```

- ▶ The above is half-hourly data so has period 48 hours and 336 hours, i.e. weekly ($336/48 = 7$)
- ▶ forecast has some special functions (notably `tbats`) for modelling multi seasonality data

18 / 27

Methods for estimating frequencies

- ▶ The most common way to estimate the frequencies in a time series is to decompose it in a *Fourier Series*
- ▶ We write:

$$y_t = \alpha + \sum_{k=1}^K [\beta_k \sin(2\pi t f_k) + \gamma_k \cos(2\pi t f_k)] + \epsilon_t$$

- ▶ Each one of the terms inside the sum is called a *harmonic*. We decompose the series into a sum of sine and cosine waves rather than with AR and MA components
- ▶ Each sine/cosine pair has its own frequency f_k . If the corresponding coefficients β_k and γ_k are large we might believe this frequency is important

19 / 27

20 / 27

Estimating frequencies via a Fourier model

- ▶ It's certainly possible to fit the model in the previous slide in JAGS, as it's just a linear regression model with clever explanatory variables
- ▶ However, it can be quite slow to fit and, if the number of frequencies K is high, or the frequencies are close together, it can struggle to converge
- ▶ More commonly, people repeatedly fit the simpler model:

$$y_t = \alpha + \beta \sin(2\pi t f_k) + \gamma \cos(2\pi t f_k) + \epsilon_t$$

for lots of different values of f_k . Then calculate the *power spectrum* as $P(f_k) = \frac{\beta^2 + \gamma^2}{2}$. Large values of the power spectrum indicate important frequencies

- ▶ It's much faster to do this outside of JAGS, using other methods, but we will stick to JAGS

21 / 27


JAGS code for a Fourier model

```
model_code = '  
model  
{  
  # Likelihood  
  for (t in 1:T) {  
    y[t] ~ dnorm(mu[t], sigma^-2)  
    mu[t] <- alpha + beta * cos(2*pi*t*f_k) +  
              gamma * sin(2*pi*t*f_k )  
  }  
  P = (pow(beta, 2) + pow(gamma, 2)) / 2  
  
  # Priors  
  alpha ~ dnorm(0, 10^-2)  
  beta ~ dnorm(0, 10^-2)  
  gamma ~ dnorm(0, 10^-2)  
  sigma ~ dunif(0, 100)  
}  
'
```

22 / 27

Example: the Lynx data

```
lynx = read.csv('../data/lynx.csv')  
plot(lynx, type = 'l')
```



23 / 27


Code to run the JAGS model repeatedly

```
periods = 5:40  
K = length(periods)  
f = 1/periods  
Power = rep(NA, K)  
  
for (k in 1:K) {  
  curr_model_data = list(y = as.vector(lynx[,2]),  
                          T = nrow(lynx),  
                          f_k = f[k],  
                          pi = pi)  
  
  model_run = jags(data = curr_model_data,  
                   parameters.to.save = "p",  
                   model.file=textConnection(model_code))  
  
  Power[k] = mean(model_run$BUGSoutput$sims.list$P)  
}
```

24 / 27

Plotting the periodogram

```
par(mfrow = c(2, 1))
plot(lynx, type = 'l')
plot(f, Power, type='l')
axis(side = 3, at = f, labels = periods)
```



25 / 27

Bayesian vs traditional frequency analysis

- ▶ For quick and dirty analysis, there is no need to run the full Bayesian model, the R function `periodogram` in the `TSA` package will do the job, or `findfrequency` in `forecast` which is even simpler
- ▶ However, the big advantage (as always with Bayes) is that we can also plot the uncertainty in the periodogram, or combine the Fourier model with other modelling ideas (e.g. ARIMA)
- ▶ There are much fancier versions of frequency models out there (e.g. Wavelets, or frequency selection models) which can also be fitted in JAGS but require a bit more time and effort
- ▶ These Fourier models work for continuous time series too

26 / 27

Summary

- ▶ We now know how to fit models for seasonal data via seasonal factors, seasonal differencing, and sARIMA models
- ▶ We can fit these using `forecast` or JAGS
- ▶ We've seen a basic Fourier model for estimating frequencies via the Bayesian periodogram

27 / 27